

BACHELOR'S THESIS

**EXPLORING REASONING
PERFORMANCE OF RISC-V
SOFTWARE MODELS IN BTOR2**

by

NADIR FEJZIĆ

submitted in partial fulfillment of the requirements
for the degree of BACHELOR OF SCIENCE
in INFORMATICS

Department of Computer Sciences
Paris Lodron University of Salzburg
Salzburg, Austria

Supervised by:

Univ.-Prof. Dipl.-Inform. Dr.-Ing. Christoph Kirsch

August 30, 2024

Abstract

When working with software programs, we want to ensure that they do not reach certain unsafe states. Various techniques exist for this purpose, such as testing. However, testing requires running the program with every possible input, which is often not practical because of the large input space. Most often we want to check whether an input exists such that our program reaches a specific state. This is known as the state reachability problem, which can be reduced to the boolean satisfiability problem and solved by SAT-solvers. The reduction can be done by encoding the sequence of machine instructions in the program as a satisfiability formula in the theory of bit-vectors, known as a Satisfiability Modulo Theory (SMT) formula. The resulting formula is the model of our program. SMT-formulae are then further reduced to boolean logic formulae by technique called bit-blasting. We can then run a SAT-solver on resulting boolean formulae, which will tell us whether the program satisfies the given constraints. Effectively, we reduced the problem of program correctness to the satisfiability problem. Important to notice is that only a finite number of instructions can be transformed this way, which makes the model checking bounded in number of executed instructions. As we reduced the program correctness problem to satisfiability problem, solving it is NP-Complete and computationally expensive. However, model generation is linear in the size of the input program. We took advantage of this property and generated many equivalent models to analyze impact of different configurations on solving performance of **btormc** model checker. In this thesis, we present a toolbox **peRIScope** for easy benchmarking of different model configurations. We also present and analyze results from benchmarking of solving time for models with different code and non-code memory granularity configurations. We found that models with smaller memory granularity were solved faster in models generated from binaries compiled with **starc** compiler. On models generated from binaries compiled with **gcc** compiler, models with larger memory granularity were solved faster.

Contents

1	Introduction	4
2	Bounded model checking	5
2.1	Correctness of software programs	5
2.2	Reduction to SAT Problem	6
2.3	Complexity control with model parameters	8
3	Model generation	9
3.1	Introduction to BTOR2	9
3.2	<code>rotor</code> - tool of choice	12
3.2.1	Model Parameters and Checks	12
4	Experiment setup	13
4.1	<code>peRIScope</code>	13
4.2	Benchmarking Workflow	14
5	Experiment results	17
5.1	Results with binaries compiled with <code>selfie</code>	18
5.2	Results with binaries compiled with <code>gcc</code>	20
6	Conclusion and further work	23

1 Introduction

One of the requirements for software programs is the correctness. Varying techniques, such as testing, are used as an attempt to ensure the correctness of written software. When testing our program we are looking for specific inputs that might be problematic and we check if our program performs as expected. In other words, we are looking for inputs for which our program reaches a certain unsafe state. However, in order to prove that the program is correct, one would have to test it with every possible input.

Running tests with every possible input is not feasible for most of practical programs because of the large input space. For programs with an 8-bit input, the number of possible inputs we need to check is $2^8 = 256$. Each additional bit doubles the number of possible inputs we would have to check, which rapidly becomes impractical. What we can do instead, is to prove that the program has the desired properties for any input. For that, we transform software programs, which are a series of machine instructions executable by a machine, to a mathematical model. In our case we transform the sequence of machine instructions to a satisfiability formula over the theory of bit-vectors and arrays of bit-vectors. Bit-vectors are signed and unsigned integers with fixed bit-length. Such formula is called a Satisfiability Modulo Theory (SMT) formula, and it uses values which are bit-vectors. In addition to instructions contained in the program, we can also add any additional constraints we want our program to uphold. Each of the used functions and predicates in an SMT-formula can be reduced to a boolean logic formula equivalent to the circuit logic of a corresponding instruction in hardware by bit-blasting. For example, addition of two 8-bit bit-vectors is reduced to a boolean formula for an 8-bit full adder. By doing this, we reduce the problem of program correctness to the problem of state reachability and then to the boolean satisfiability problem. We can now run a SAT-solver on the resulting boolean formula, which will tell us whether our program satisfies the desired constraints. This reduction is limited to a finite number of machine instructions, which makes the model checking bounded in number of executed instructions.

Since the model is a reduction of the correctness problem to a satisfiability problem, solving the model means solving the satisfiability problem. The satisfiability problem is, however, NP-Complete. This means that solving of such models is computationally expensive. On the other hand, model generation can be computed linearly in the size of the program (number of instructions).

Fast model generation can be taken advantage of and we can generate multiple equivalent models with different configurations in order to influence the complexity of resulting boolean formulae. To better illustrate this, we can analyze, for example, the memory access circuitry in machines and, therefore, the boolean logic formulae for memory access in generated models.

Depending on the size of each addressable memory block (memory granularity), we need different number of addresses to access the whole memory. Machines have logic circuits for memory address decoding combined with logic circuits for accessing specific memory lines. These logic circuits are encoded in our models as well. With smaller memory granularity, we need more addresses to access the whole memory. Opposite is the case for larger memory granularity. So the complexity of logic circuits for memory access is directly influenced by the memory granularity. Larger memory granularity is better in this regard, but worse in others. In general, extracting values from memory

that differ in number of bits from the used memory granularity increases complexity. We might need to use bit-shifting, bit-masking and even multiple loads or stores in order to manipulate the desired value. This logic for reading and writing of values from and to memory is encoded in our models as well. All this influences complexity in models of programs with frequent access to values in memory smaller or larger than the configured memory granularity. By changing the memory granularity we shift the complexity of models between the formulae for memory address decoding and memory access logic, and formulae for logic of reading and writing of values from and to memory.

We can tweak many parameters of models to analyze how these shifts in complexity affect the performance of `btormc` model checker. In this thesis we're particularly interested in solving performance when using different combinations of memory granularity for code and non-code memory segments. We present a toolbox `peRIScope` that orchestrates the benchmarking of model solving. This includes generation of models with different configurations, running the model checker on generated models and tracking of the time taken for solving.

2 Bounded model checking

2.1 Correctness of software programs

Software programs are written with a particular goal in mind. We can create a specification that describes what the goal of the program is. If the program does what it's supposed to do as defined by the specification, we say that the program is correct. Depending on the program and our requirements, we can define different specifications for different types of correctness. For example, we can check whether the program logically does what it's supposed to. This would be logical correctness. Another example is we can check whether the program performs unsafe operations, such as accessing memory out of bounds. This would be safety correctness and so on.

In the context of this thesis, correctness of the program is defined as:

- Program does not terminate with a bad exit code
- Program terminates with a good exit code
- Program does not perform division by zero
- Program does not inhibit division overflow
- Program does not access any invalid memory address
- Program does not inhibit segmentation fault

However, testing for each of these properties for any input is problematic. If our program has a single 8-bit integer as its input, we would have to test it with 2^8 different inputs. For each bit added to the input, regardless if as an additional input or as an additional bit to the existing input, we double the number of tests. Testing with such large input space is not feasible for most programs.

A better approach is to prove that our program is correct for any given input. In other words, we want to prove that given an input, the program reaches a specific state [5]. What we're interested in, is whether an input for the given program exists, such




Figure 1: Schematic of full adder implemented with two XOR gates, two AND gates, one OR gate.

that a machine reaches a certain state while it executes that program. This is known as the state reachability problem [1]. In this thesis we further reduce the state reachability problem to the boolean satisfiability problem. The formal definition of the problem we want to solve is called model. The reduction is done in such a way, that our model is satisfiable if and only if the specified state is reachable. This means that, if our model is satisfiable, our program upholds the desired constraints. Analog if we do not want the program to reach a specific state, then satisfiable model means that our program violates some constraints [3].

2.2 Reduction to SAT Problem

A program is a series of instructions that can be executed by a particular machine. The machine fetches instructions, decodes and executes them. Instructions can have various semantics such as arithmetic operations, memory access, branching, and so on. By executing the instructions, machine changes its state. Minimal amount of state that is changed for each instruction is the update of the program counter. Apart from program counter, the machine can modify its registers and main memory as well [7].

Instructions in machines are implemented using logical gates, such as AND, OR, NOT and other gates. Logical gates are devices that perform boolean functions and are implementations of boolean formulae. In other words, we can model the logical gates of the machine using boolean formulae [14]. The logical circuit of one instruction becomes a boolean formula in our model. An example of a 1-bit full adder schematic can be seen in figure 1. The schematic shows circuit logic for arithmetic addition with two 1-bit inputs A and B , one 1-bit value from previous addition called carry in C_{in} and it computes the 1-bit sum S and 1-bit carry out C_{out} value in case of an overflow. If we

have a simple RISC-V program containing the following series of instructions:

```
addi t0 , zero , 4
addi t1 , zero , 2
mul  t0 , t0 , t1
```

then we can model the program with the following operations:

```
t0_1 = 0      + 4      && pc_1 = pc_0 + 4 &&
t1_1 = 0      + 2      && pc_2 = pc_1 + 4 &&
t0_2 = t0_1 * t1_1 && pc_3 = pc_2 + 4
```

Listing 1: Example model

So the small RISC-V program can be reduced to the series of operations in listing 1. What we encoded is a satisfiability modulo theory formula (SMT-Formula). That is, we encoded the program as a series of operations in theory of bit-vectors. This model is satisfiable if and only if there exists an assignment of values to variables such that the formula is true. Variables in an SMT-Formula are bit-vectors and arrays of bit-vectors, hence the theory of bit-vectors. Bit-vectors represent signed and unsigned integers of arbitrary but fixed bit-length [4], and bit vector arrays which contain multiple values (either bit-vectors or arrays), indexable by a bit-vector [13]. Since the model we encoded operates on bit-vectors and arrays of bit-vectors and not on boolean values, it is not boolean logic. It is an intermediate step that must be further reduced to boolean formula.

If we recall the full adder schematic in figure 1, we can see that the logic gates correspond to the following boolean formulae:

$$S = A \oplus B \oplus C_{in}$$

$$C_{out} = (A \wedge B) \vee (C_{in} \wedge (A \oplus B))$$

This is a series of boolean formulae that describe the behaviour of the 1-bit full adder. For two bit full adder, we would have two sum outputs S_1, S_2 , one initial carry in C_{in1} , and the carry out of the first sum C_{out1} would be the carry in to the second sum $C_{out1} = C_{in2}$. The same process applies for full adder of values with arbitrary bit length. The series of operations in listing 1 can be encoded in a similar way. This technique of encoding semantics of arithmetic and bitwise operations is called bit-blasting [2]. We can add more constraints to the model that do not come from the program, but rather specify whether the machine reaches a particular state. For example, we can add constraints that check whether a certain register of the machine (variable) contains the value 0 and we perform a division with that register.

By doing so, we formulated a satisfiability problem, and by solving it we prove that either our program does not perform division by zero for any given input, or we get an example input for which the program performs division by zero. In particular, by doing this, we reduced the state reachability problem to the boolean satisfiability problem. Important observation is that the resulting boolean formula is finite and can model a finite number of executed instructions, which makes our models bounded in the number of executed instructions. This is the essence of bounded model checking.

2.3 Complexity control with model parameters

There are multiple possibilities when generating model and certain trade-offs can be made. In our models of the machine, bit-vectors are used for values and arrays of bit-vectors for main memory and registers. In particular, we used one-dimensional arrays which contain bit-vectors of some fixed bit-length. When modelling M -bit memory with an array that contains N -bit bit-vectors we need an address space of $2^{(M-N)}$ entries. More precisely, the array must be indexed with a bit-vector of bit length $(M - N)$. Size of bit-vectors stored in the array represent the memory granularity - the size of each memory block our machine can access. By changing the memory granularity we can directly impact the size of the array.

When accessing memory each memory block is addressable by an address, which is a bit-vector in our case. Logical circuit exists for this purpose in machine as well. Naively implemented, this could be a chain of something similar to if-else statements. In reality machines use a logical circuits such as decoders, multiplexers and demultiplexers. Decoders have an N -bit address as input which is decoded and used to activate one of 2^N possible outputs. Multiplexer has 2^N inputs, one of which is selected based on the output produced by the decoder given an N -bit address. Analog, demultiplexer has one input which is routed to one of 2^N outputs based on the address [9]. Circuit logic of these components can be modelled as boolean formulae as well, effectively modelling memory access of the machine. Number of inputs for multiplexer (or outputs for demultiplexer) multiplied by the size of each input is the size of modelled memory. The number of outputs of address decoder must match the number of inputs/outputs of (de-)multiplexer. Size of each input/output corresponds to memory granularity, and so the memory granularity directly impacts the complexity of resulting circuit logic for memory access and therefore the complexity of the model. Smaller memory granularity requires more lines for input/output and larger memory addresses, which increases complexity in all of the mentioned components.

On the other hand, if we choose larger memory granularity, we can access the whole memory with fewer addresses. Consequently, this reduces the complexity of memory access circuits and therefore the complexity in corresponding part of our model. However, it is not always the case that we need to access values in memory that have the same size as our chosen granularity. When that's not the case, we have to perform bit-shifting and bit-masking operations in order to load the correct value. For example, if we want to load a 32-bit value from memory with 8-bit granularity, we will have to access memory four times and perform bit-shifting to store bits of value at the correct offset. Smaller granularity in this case results in four memory accesses through large circuit and four bit-shifting operations. If we, however, want to load a 32-bit value from memory with 64-bit granularity, complexity varies greatly depending on the alignment of the value in memory. As we can see in figure 2, when loading a value that is fully contained inside of a 64-bit block, a single load with bit shifting and/or bit masking will suffice. But if the value is split between two 64-bit blocks, we need two loads, both bit shifting and bit masking and finally combining the two values. Also, determining whether the value is split between two blocks requires division operations to be performed, increasing the number of operations on each access.

In our models, we can control the memory granularity and therefore influence the complexity of models in various parts. In particular, we can control two different parts



in BTOR2 can define operations such as addition, multiplication etc., or they can be variable declaration, memory access, constraints and so on [13].

In listing 2 we can see an example of a model in BTOR2. In first two lines we define two sorts, one is a 1-bit bit-vector with sort id 1, and the other is a 32-bit bit-vector with sort id 2. In the following line we have a node with node id 3 that defines an input variable `turn` of sort with sort id 1, meaning that `turn` is a 1-bit bit-vector. In the next line we define a constant 0 of sort with sort id 2. In next two lines we define two 32-bit bit-vector state variables `a` and `b`, and then initialize them with previously defined constant 0. We then define a constant 1, which we use to increment the state variables. We can observe that most of the keywords first reference a sort id. The referenced sort id is the type of the node, and the node itself evaluates to a certain value that we can reference by the node's id. In nodes 12 and 13, we use an if-then-else construct, where we evaluate to a certain value based on the condition. For state variable `a` we check whether the value of node 3 is 1 (*true*), and if it is, the node 12 evaluates to value of node 5, otherwise it evaluates to the value of node 10. For state variable `b` we check whether the value of node 3 is false (-3 negates the value of node 3), and if so we use the value of node 6, otherwise value of node 11. Since nodes 5 and 6 are definitions of state variables, using them means using whatever the current value of the variable is. In nodes 14 and 15 we define the next function, which is the function that updates the state variables, and these functions use the previously defined if-then-else nodes for the update. In the next few lines we define a 32-bit bit-vector constant 3, we have two nodes that check whether our state variables have value 3. If both of them have value 3, the node 19 evaluates to 1 (*true*), and in node 20 we define a bad state. The node 20 simply means, if the value of node 19 is 1, then this bad state is reached. In listing 3 we can see the corresponding C code for the BTOR2 model.

```

1 sort bitvec 1
2 sort bitvec 32
3 input 1 turn
4 zero 2
5 state 2 a
6 state 2 b
7 init 2 5 4
8 init 2 6 4
9 one 2
10 add 2 5 9
11 add 2 6 9
12 ite 2 3 5 10
13 ite 2 -3 6 11
14 next 2 5 12
15 next 2 6 13
16 constd 2 3
17 eq 1 5 16
18 eq 1 6 16
19 and 1 17 18
20 bad 19

```

Listing 2: Example BTOR2 model for C code in listing 3

```

int main(void) {
    bool turn;
    unsigned int a = 0;
    unsigned int b = 0;

    while (1) {
        turn = read_bool();
        assert(!(a == 3 && b == 3));

        if (turn)
            a += 1;
        else
            b += 1;
    }
}

```

Listing 3: C code we generate a model from

The BTOR2 format encodes the SMT-Formula. To solve the formula, we use a set of tools that accompany the BTOR2 format available as part of the Boolector project [12]. In particular, we use the reference implementation of a bounded model checker **btormc**. The bit-blasting of SMT-Formula is performed by the **btormc** before solving it. We run the model checker by providing it the file containing BTOR2 model, and optionally the upper bound. Upper bound option sets the maximal number of instructions that can be checked. If the model checker terminates with no output, that means that our program does not satisfy the given constraints in the number of instructions we chose as the upper bound. On the other hand, if the program satisfies the constraints, a witness format is produced. We can see an example output of witness format in listing 4.

```

sat
b0
@0
0 1 turn@0
@1
0 0 turn@1
@2
0 0 turn@2
@3
0 0 turn@3
@4
0 1 turn@4
@5
0 1 turn@5
@6
0 0 turn@6
.

```

Listing 4: BTOR2 Witness Format for model in listing 2

The `sat` indicates that the model is satisfiable, `b0` is the satisfiable property with `b0` for 0-th bad property, that is the first bad property that appears in the model. `@0` through `@6` indicate the assignments of the inputs in frames 0 to 6, e.g., e.g., in frame 0 the 0-th input (first input defined in the BTOR2 model) `turn = 1` and so on. Frame 6 is the last frame, where the bad property `b0` is satisfied.

3.2 rotor - tool of choice

Generating BTOR2 models for programs is not done by hand. That would be tedious and error prone, as models tend to be very large counting tens of thousands of lines. Instead, we use a tool capable of generating BTOR2 models. Specifically, we use `rotor`, which is a self-translating modeling engine based on `selfie` that translates full RISC-V code including all of `selfie` and itself to BTOR2 and SMT-LIB formulae that are satisfiable if and only if there is input to the code such that the code exits with non-zero exit codes, performs division by zero, or accesses memory outside of memory segments. Rotor also generates models that enable RISC-V code synthesis [11].

3.2.1 Model Parameters and Checks

`rotor` accepts multiple arguments and parameters that can be used to adjust produced models. By default multiple checks are enabled in `rotor`, and we have to explicitly disable them using the following command line arguments:

- `-Pnobadexitcode` - Disables check for bad exit code. Without this option, the model checker will report satisfiable if the program can terminate with the chosen bad exit code.
- `-Pgoodexitcode` - Enables check for good exit code. With this enabled, the model is satisfiable if the program can exit with some exit code other than the chosen good exit code.
- `-Pnoexitcodes` - Disables checking whether multiple cores exit with the same exit code. If not present, model is satisfiable if different cores can exit with different exit codes.
- `-Pnodivisionbyzero` - Disables division by zero check, which checks whether an input for the program exists such that the program performs division by zero.
- `-Pnodivisionoverflow` - Disables check for division (or remainder) overflow. Without this option, model is satisfiable if an input exists such that program performs division with an overflow.
- `-Pnosegfaults` - Disables check for segfaults, which checks whether the program accesses memory outside of defined memory segments.

Apart from these checks, we can also modify some parameters of the model. Models generated for benchmarking in this thesis are generated with following parameters:

- `-bytestoread 1` - Input of the program is 1 byte
- `-cores 1` - Machine has a single core
- `-virtualaddressspace 32` - 32-bit virtual address space

- `-codewordsize X` - Granularity of the code memory segment, where `X` is the size of the code memory segment in bits
- `-memorywordsize X` - Granularity of the non-code memory segments (data, stack and heap), where `X` is the size of the memory segment in bits.
- `-heapallowance 4096` - Maximum heap size in bytes
- `-stackallowance 4096` - Maximum stack size in bytes

As we can see, the memory granularity of code and non-code memory segments can be independently configured. This was used to generate models with different combinations of code and non-code memory granularity. We can generate many models rather quickly, because the generation of models is linear to the size of the input program. `rotor` accepts either `C*` source code, binary compiled by `starc` or binary compiled by `gcc` as inputs. In case of using binaries as input, they must be compiled for the RISC-V architecture.

4 Experiment setup

Running the experiment consists of multiple steps. This includes obtaining binaries compiled for RISC-V architecture and generating models in BTOR2 format from those binaries. After that, `btormc` is run on each model multiple times, measuring and tracking the time it took to solve the model. We also want to confirm that models are correct and are satisfiable for the right reason. With the goal to facilitate this process, we introduce `peRIScope`.

4.1 peRIScope

`peRIScope` is a benchmarking toolbox designed to make the process of parameterized model generation and benchmarking of model solving easier [6]. The tool also supports parsing of BTOR2 witness format in order to provide more information, such as which property was satisfied, and transitions of inputs through frames that satisfy said property. `peRIScope` accepts a configuration file in either `YAML` or `JSON` format, which can be used to specify the parameters for models, timeout for maximum duration of model solving, list of files for filtering, and flags for `btormc`. The `peRIScope` performs multiple runs, each time invoking `rotor` with parameters of that run which results in multiple models being generated in a (specific) subdirectory inside `selfie`. Model generation is fast, so we can generate models for all example files without filtering. After the models are generated, `btormc` is run on each model multiple times using `hyperfine` [15], measuring the time it took each time. In the example configuration file shown in listing 5 we tell the `peRIScope` to generate models with different code memory granularity and to benchmark only models that match any of the entries in the `files` list. The `timeout` parameter sets the maximum number of seconds solving is allowed to take.

```

timeout: 300
files:
  - "division-by-zero-3-35-rotorized.btor2"
  - "invalid-memory-access-fail-2-35-rotorized.btor2"

runs:
  8-bit-codeword-size: "0 -codewordsize 8"
  16-bit-codeword-size: "0 -codewordsize 16"
  32-bit-codeword-size: "0 -codewordsize 32"
  64-bit-codeword-size: "0 -codewordsize 64"

```

Listing 5: Example peRIScope configuration file in YAML format

4.2 Benchmarking Workflow

The peRIScope can be compiled to a binary and at the time of this writing supports two commands. One command is **parse-witness** which takes witness format of **btormc** as input. The witness format is parsed, and the flow of inputs and states through frames is printed to **stdout** in a condensed format, as can be seen in listing 6. This command has an optional parameter which is used to provide the path to BTOR2 model file for which the witness format was generated. peRIScope reads the BTOR2 file and extracts more information about the model such as the names of properties that were satisfied. This enables peRIScope to display additional information in the result of witness parsing and evaluation, like in listing 7.

```

Satisfied properties in 77 steps:
  Bad at 7

Inputs flow:
States flow:
  input-buffer:
    #0: 48 ([00000000] -> 00110000)
    -> #1: end

```

Listing 6: Example output of peRIScope parse-witness command.

```

Satisfied properties in 77 steps:
  Bad at 7 named 'core-0-division-by-zero' with nid: 37027

Inputs flow:
States flow:
  input-buffer:
    #0: 48 ([00000000] -> 00110000)
    -> #1: end

```

Listing 7: Example output of peRIScope parse-witness command with provided BTOR2 file for context.

The output of peRIScope command **parse-witness** describes that the given properties were satisfied in some number of frames. Additionally, it shows that the **input-buffer**




Figure 3: Plot of results for two BTOR2 models with log scale.

variable, modelled as `state` in BTOR2, is assigned value 48 at frame 0, and it did not change until the very last frame.

The second command of `peRIScope` is the `bench` command. This command orchestrates the execution of multiple different commands. First, it runs the `rotor` generating models with specific configuration. Then it runs `hyperfine` executing `btormc` on each model, collecting benchmarking statistics. `btormc` may run very long for certain models, so if the `timeout` option is set in the used configuration, we wrap the invocation in the `timeout` command of GNU `coreutils` [8]. We mark the result as `Failed` if we use `timeout` and the model solving does not terminate within the specified time. Otherwise, witness format is generated, parsed alongside the provided BTOR2 model for context. The result of the whole run is then stored in a results file, example of which can be seen in listing 8.

The results are stored in JSON format, with one entry per BTOR2 model file. The `Success` key indicates that `btormc` terminated, and list of satisfied properties is presented. Further, we store the number of frames it took to satisfy the properties and the results of `hyperfine` benchmark runs, containing statistics as well as specific timings and exit codes of each run. We also store the word count in BTOR2 model file, and also word count of `btormc` dump for the same model file after it processed it. This is useful for tracking the correlation between model size and the time it took solving it.

Accompanying the `peRIScope` binary, we also have a script written in Python. The script takes the results file as input and generates plots for easier visual inspection of the results. The script supports different types of plots, such as box plots and histograms. Linear and log scaling of plots are available. Figure 3 is one such plot.

The models can be generated from C* source code. `rotor` handles this by first compiling the source code with `starc` and then generating the BTOR2 model from the resulting binary. The models can also be generated from binaries directly. `rotor` has a flag `-l` for loading a binary file. Binaries compiled with either `starc` or `gcc` are supported. The `starc` compiles binaries for RISC-V architecture out of the box. In case of `gcc` we either have to compile the binaries on a RISC-V machine, or we have to use the RISC-V toolchain for cross-compilation [10]. In this thesis, we compiled binaries and generated models from the same set of C* source code files with both `starc` and `gcc`. In particular, we used the files found in `examples/symbolic/` directory in the `selfie` repository [11].

The benchmarks were run on a machine with following specifications:

- Huawei Matebook 16 (model: CREM-WFD9)

- CPU: AMD Ryzen 7 5800H (16-core) clocked at 4.46 GHz
- Graphics: AMD Radeon Vega Mobile Series (Integrated)
- Memory: 16GB LPDDR4-4167

```
{
  "division-by-zero-3-35-rotorized.btor2": {
    "Success": {
      "props": [
        {
          "kind": "Bad",
          "name": "core-0-division-by-zero",
          "node": 37027,
          "idx": 7
        }
      ],
      "steps": 77,
      "hyperfine": {
        "results": [
          {
            "command": "btormc division-by-zero-3-35-rotorized.btor2 -kmax 200",
            "mean": 0.3718056574400001,
            "stddev": 0.008814608910627868,
            "median": 0.37023892844,
            "user": 0.33166193999999993,
            "system": 0.0399969,
            "min": 0.36236449744,
            "max": 0.38613962244,
            "times": [
              0.37023892844,
              // ...
            ],
            "exit_codes": [
              0,
              // ...
            ]
          }
        ]
      },
      "wc_raw": 273051,
      "wc_btormc_dump": 93238
    }
  }
}
```

Listing 8: Example of peRIScope results file

5 Experiment results

With the setup described in previous section, we generated BTOR2 models with following examples:

- Simple assignment - satisfiable on specific input manipulated with basic arithmetic.
- Division by zero - satisfiable if division by zero is performed.
- Invalid memory access fail - satisfiable if invalid memory is accessed, such as memory at an address outside of the virtual address space.
- Memory access fail - satisfiable if bad memory access is performed, such as unaligned memory access.
- Simple if without else - satisfiable when the input satisfies the condition of an if statement without an else block.
- Simple if-else - satisfiable when the input satisfies the condition of an if statement.
- Simple if-else reverse - satisfiable if the input does not satisfy condition of an if statement so that the body of the `else` block is executed.
- Nested if-else - satisfiable when body of an if statement inside of another if statement is executed.
- Nested if-else reverse - satisfiable when body of an if statement inside of an else block is executed.
- Recursive ackermann - satisfiable if ackermann function is computed with a specific input.
- Recursive fibonacci - satisfiable if fibonacci function is computed with a specific input.
- Recursive factorial fail - satisfiable if factorial function is computed with a specific input.
- Nested recursion fail - satisfiable if nested recursion is performed with a specific input.
- Return from loop - satisfiable if a loop exits with a return statement, which happens with digit as the input.
- Simple increasing loop - satisfiable if a loop is run exact number of times, depending on the input value used in the loop condition which is incremented in each iteration.
- Simple decreasing loop - satisfiable if a loop is run exact number of times, depending on the input value used in the loop condition which is decremented in each iteration.
- Three level nested loop fail - satisfiable if three nested loops that use input in their condition run specific number of times.
- Two level nested loop - satisfiable if two nested loops that use input in their condition run specific number of times.

Each of the examples is contained in separate source code file. We compiled each of the files both with **starc** and **gcc** compilers, and generated BTOR2 models with different combinations of code-word and memory-word sizes. The benchmark results of bounded model checking (with 200 as the upper bound) for each model are presented in the following two sections.

5.1 Results with binaries compiled with selfie




Figure 4: Selfie compiled binaries - codeword size comparison

The time it took to solve BTOR2 models generated from binaries compiled with **starc** compiler in **selfie** are presented in figures 4 and 5. The solving terminated within 300 seconds for 9 models configured with code-word size of 8-bits. In case of models with 32-bit code-word size, solving terminated for 10 such models, and for 11 models with 64-bit code-word size. As for the configuration with 16-bit code-word size, solving terminated within 300 seconds for 12 models. This configuration terminated approximately in the average duration of all other models. However, solving terminated for more models with this configuration than with any other. This indicates that time complexity grows slower for models with 16-bit code-word size generated from **starc** compiled binaries. This is somewhat surprising, as **starc** compiled binaries contain only 32-bit RISC-V instructions.

We proceeded with the 16-bit code-word size configuration, and generated models with varying memory-word sizes. Comparison of solving time for these models in figures 6 and 7 show that the configuration with 8-bit memory-word size terminated within 300 seconds for 15 models, 16-bit memory-word size for 13 models, 32-bit memory-word as well as 64-bit memory-word size configurations terminated for 13 models. The configuration with 8-bit memory-word size was solved in shortest amount of time for 6




Figure 5: Selfie compiled binaries - codeword size comparison (Log scale)




Figure 6: Selfie compiled binaries - fixed 16-bit codeword with varying memory word size comparison

models. 16-bit memory-word size configuration was solved in shortest amount of time for 4 models, 32-bit memory-word size configuration for 3 models and 64-bit memory-




Figure 7: Selfie compiled binaries - fixed 16-bit codeword with varying memory word size comparison (Log scale)

word size configuration for 2 models. Overall models generated from **starc**-compiled binaries with smaller memory granularity configurations terminated in shorter amount of times.

5.2 Results with binaries compiled with gcc

The times it took for **btormc** to terminate for models with different codeword-size configurations from binaries compiled with **gcc** are presented in figures 8 and 9. The **btormc** terminated within 300 seconds for roughly half of the generated models. This was expected, because the models generated from binaries compiled with **gcc** are much bigger. The results can be better analyzed in the figure 9 with logarithmically scaled time axis. Relations between 8-bit and 16-bit codeword size configurations are similar to those between 32-bit and 64-bit codeword size. The models with 8-bit codeword size configurations were solved faster than those with 16-bit configurations in all but two cases. Similarly for 32-bit and 64-bit codeword size configurations. However, in those cases where the models with smaller codeword memory granularity were solved faster, the difference in time is insignificant, approximately half of a second. But in cases where models with larger codeword size were solved faster, the time difference is bigger, around a second between 64-bit and 32-bit configurations and around 30-40 seconds between 16-bit and 8-bit configurations. Overall, the shortest time it took to solve the models was for configurations with 64-bit codeword size.

As the models with 64-bit codeword size configuration were the ones that were overall solved the fastest, we proceeded with this configuration, generated models with varying memory-word sizes and obtained results shown in figures 10 and 11. **btormc** terminated

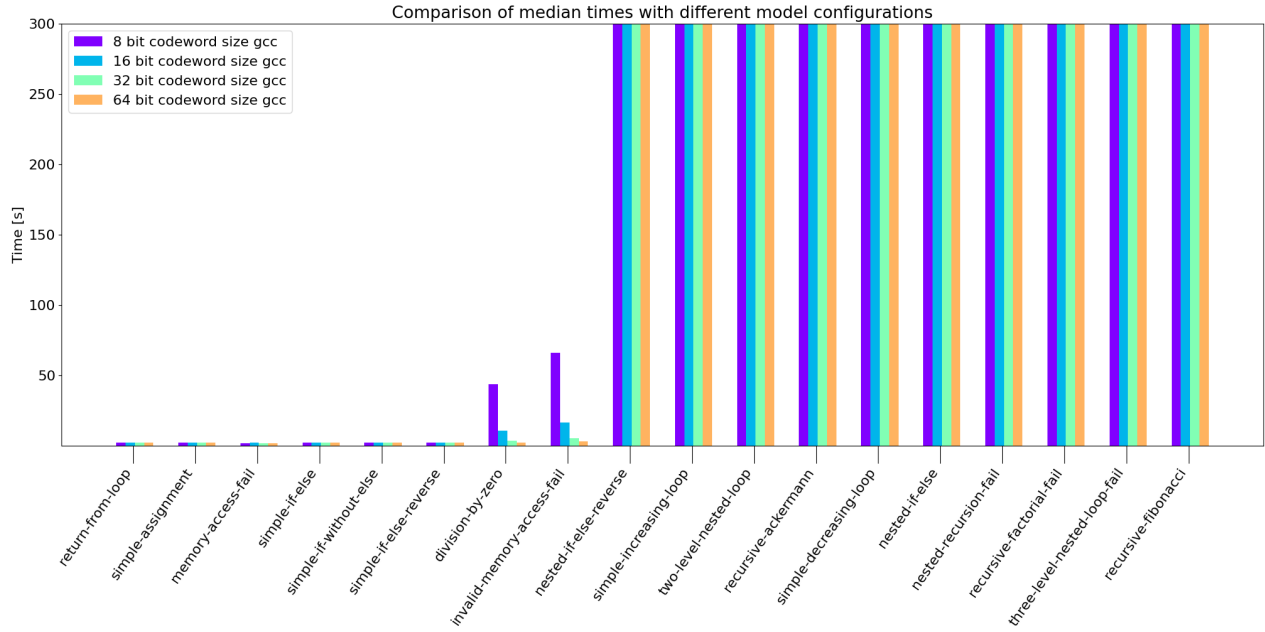


Figure 8: gcc compiled binaries - codeword size comparison




Figure 9: gcc compiled binaries - codeword size comparison (Log scale)




Figure 10: gcc compiled binaries - fixed 64-bit codeword, varying memory word size comparison




Figure 11: gcc compiled binaries - fixed 64-bit codeword, varying memory word size comparison (Log scale)

for the same number of models. We can clearly see that the models configured with 64-bits for both codeword and memoryword sizes were solved in the shortest amount of time. The models with 32-bit memory-word size come closely in second place. In case of 8-bit and 16-bit memory-word size configurations, 8-bit configuration was solved faster in all but 2 cases. Overall, we observe that in case of `gcc` compiled binaries larger memory granularity seems to be more beneficial.

6 Conclusion and further work

In this thesis, we discussed the importance and proving of program correctness. We explained shortly, how to model the problem of proving correctness by reducing the state reachability problem to the boolean satisfiability problem. We showed how to generate SMT-Formula from a program, and then further reduce it to a boolean formula by bit-blasting. Two new tools were presented, `rotor` used for BTOR2 model generation, and `peRIScope` for benchmarking orchestration of bounded model checking of generated models. Finally, we ran the BTOR Model Checker `btormc`, measured the time it took to solve the models generated from binaries compiled with `starc` and `gcc`, and briefly discussed the results we observed. The results indicate that solving of models generated from binaries compiled by `starc` is faster for configurations with smaller code and non-code memory granularity. In case of `gcc` compiled binaries, solving was faster for models with largest code and non-code memory granularity. This might be due to multiple reasons, such as the size of models, but also potential optimizations implemented by `gcc` for 64-bit RISC-V architecture. Future experiments could investigate the correlation between compiler optimizations and the solving time.

The generated model configurations are not exhaustive. `rotor` supports more options, such as specifying the number of cores, how many bytes of input should be read and so on, and it can be extended to support even more options. More exhaustive benchmarking of every possible combination would provide more insight into the reasoning performance of RISC-V Software Models in BTOR2.

References

- [1] Parosh Aziz Abdulla, Per Bjesse, and Niklas Eén. “Symbolic reachability analysis based on SAT-solvers”. In: *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer. 2000, pp. 411–425.
- [2] Clark Barrett et al. “Satisfiability modulo theories”. In: *Handbook of satisfiability*. IOS Press, 2021, pp. 1267–1329.
- [3] Robert Brummayer, Armin Biere, and Florian Lonsing. “BTOR: bit-precise modelling of word-level problems for model checking”. In: *Proceedings of the joint workshops of the 6th international workshop on satisfiability modulo theories and 1st international workshop on bit-precise reasoning*. 2008, pp. 33–38.
- [4] Leonardo De Moura and Nikolaj Bjørner. “Satisfiability modulo theories: An appetizer”. In: *Brazilian Symposium on Formal Methods*. Springer. 2009, pp. 23–36.
- [5] Douglas D Dunlop and Victor R Basili. “A comparative analysis of functional correctness”. In: *ACM Computing Surveys (CSUR)* 14.2 (1982), pp. 229–244.
- [6] Nadir Fejzić. *peRIScope*. URL: <https://github.com/nfejzic/periscope>.
- [7] Michael Flynn. “Computer architecture”. In: *Wiley Encyclopedia of Computer Science and Engineering* (2007).
- [8] *GNU Coreutils*. Version 9.0. Free Software Foundation, August 25, 2024. URL: <https://www.gnu.org/software/coreutils/>.
- [9] Paul Horowitz, Winfield Hill, and Ian Robinson. *The art of electronics*. Vol. 2. Cambridge university press Cambridge, 1989.
- [10] RISC-V International. *RISC-V GNU Compiler Toolchain*. URL: <https://github.com/riscv-collab/riscv-gnu-toolchain>.
- [11] Christoph Kirsch. *selfie*. August 2024. URL: <https://github.com/cksystemsteaching/selfie>.
- [12] Aina Niemetz, Mathias Preiner, and Armin Biere. “Boolector 2.0”. In: *J. Satisf. Boolean Model. Comput.* 9.1 (2014), pp. 53–58. DOI: 10.3233/sat190101. URL: <https://doi.org/10.3233/sat190101>.
- [13] Aina Niemetz et al. “Btor2, btormc and boolector 3.0”. In: *International Conference on Computer Aided Verification*. Springer. 2018, pp. 587–595.
- [14] Daniel Page. *A practical introduction to computer architecture*. Springer Science & Business Media, 2009.
- [15] David Peter. *hyperfine*. Version 1.16.1. March 2023. URL: <https://github.com/sharkdp/hyperfine>.

- [16] Andrew Waterman et al. “The RISC-V instruction set manual, volume I: User-level ISA, version 2.0”. In: *EECS Department, University of California, Berkeley, Tech. Rep. UCB/EECS-2014-54* (2014), p. 4.

Statement of Authentication

I hereby declare that I have written the present thesis independently, without assistance from external parties and without use of other resources than those indicated. The ideas taken directly or indirectly from external sources (including electronic sources) are duly acknowledged in the text. The material, either in full or in part, has not been previously submitted for grading at this or any other academic institution.

Salzburg, August 30, 2024

Nadir Fejzić